

Gait Anomaly Detection via Clinically Motivated IMU Perturbations: A Comparative Study of Statistical, Ensemble, and Deep Learning Approaches under Cross-Subject Generalisation

Mouna Ben Amor

Higher Institute of Computer Science and Multimedia of Gabes
mouna.benamor@isimg.tn

Abstract—Wearable inertial measurement units (IMUs) enable continuous, unobtrusive monitoring of human gait, offering clinical value in fall prevention, rehabilitation tracking, and neurological disease management. We frame gait anomaly detection as an unsupervised learning problem and evaluate three detectors—a statistical threshold model, an Isolation Forest, and an LSTM Autoencoder—on four clinically motivated perturbation types injected into normal walking acceleration windows from the UCI HAR Dataset [12]. Anomaly types model progressive fatigue drift, sensor placement noise, compensatory gait asymmetry, and near-fall impulsive events. We evaluate under two protocols: a pooled random split and a cross-subject protocol that withholds nine subjects entirely from training. The LSTM Autoencoder achieves the best clinical operating point—perfect precision (1.000) with MCC of 0.711 on the pooled split and 0.646 cross-subject—making it the most deployable detector in low-false-alarm settings. SHAP attribution identifies spectral energy features (`std_fft_magnitude`, `mean_fft_magnitude`) as primary drivers of anomaly detection, with per-type attribution revealing mechanistically distinct feature signatures for each clinical condition. Notably, models perform comparably or better under cross-subject evaluation than pooled, a finding we attribute to anomaly-type distribution variance between splits. Code is publicly available at <https://github.com/itsMouna/gait-anomaly-detector>.

Index Terms—gait analysis, anomaly detection, IMU, wearable sensing, isolation forest, LSTM autoencoder, SHAP, fall detection, human activity recognition, cross-subject generalisation

I. INTRODUCTION

Falls are the leading cause of injury-related death among adults over 65, with an estimated 684,000 fatal falls occurring globally each year [1]. Continuous gait monitoring via wearable IMUs offers a promising avenue for early detection of pathological gait changes that precede falls [2].

Existing systems typically require labelled pathological recordings, limiting deployability in real-world settings where such data is unavailable [3]. Unsupervised detectors—trained only on normal gait—are more practical but have received less rigorous evaluation with respect to clinically meaningful anomaly types and cross-subject generalisation.

This paper makes four contributions:

- 1) A **clinically grounded anomaly taxonomy** mapping four injection modes to documented gait pathologies with supporting literature.
- 2) A **dual evaluation protocol** contrasting pooled splits with strict cross-subject evaluation on nine held-out individuals.
- 3) **SHAP attribution** at both global and per-anomaly-type resolution, exposing mechanistically distinct feature signatures.
- 4) A **clinical operating-point analysis** arguing that perfect precision is more deployment-relevant than maximum AUC in gait monitoring scenarios.

II. RELATED WORK

A. Wearable Gait Analysis

IMU-based gait analysis has been applied to Parkinson's disease monitoring [5], post-stroke rehabilitation [6], and fall risk assessment [4]. Spectral features computed from accelerometer windows have shown strong discriminative power for activity classification [2].

B. Anomaly Detection Methods

Isolation Forest [9] isolates outliers by recursively partitioning the feature space; anomalies require fewer cuts and receive higher anomaly scores. LSTM-based autoencoders [10] learn to reconstruct normal windows; anomalous windows produce elevated reconstruction error. Statistical threshold detectors provide interpretable clinical baselines [7].

C. Explainability in Clinical AI

SHAP [11] decomposes predictions into per-feature Shapley contributions. In clinical settings, attribution is essential for regulatory compliance and clinician trust. TreeSHAP [11] provides exact SHAP values for tree-based models at polynomial cost.

III. DATASET

We use the UCI HAR Dataset [12], collected from 30 volunteers (ages 19–48) wearing a Samsung Galaxy S II at the waist. Body acceleration was recorded at 50 Hz and pre-segmented into 128-sample (2.56 s) windows. We retain WALKING windows only (label=1), yielding 1,226 windows across all subjects. The magnitude signal $m_t = \sqrt{x_t^2 + y_t^2 + z_t^2}$ is used throughout.

IV. METHOD

A. Clinically Motivated Anomaly Injection

We inject four perturbation types into 15% of test windows (Table I). Each type is grounded in a documented gait pathology.

TABLE I
CLINICAL ANOMALY TAXONOMY

Type	Perturbation	Clinical analogue
Fatigue drift	+ linspace(0, 0.8)	Muscle fatigue [4]
Sensor noise	+ $\mathcal{N}(0, 0.5^2)$	IMU attachment error [8]
Gait asymmetry	$\times 0.6$	Post-stroke hemiparesis [6]
Near-fall event	5-sample spike +2.0	Stumble impulse [7]

B. Feature Extraction

We extract seven features per window:

$$\mathbf{f}(w) = [\mu_w, \sigma_w, \overline{w^2}, \max w, \min w, |\hat{w}|, \sigma_{|\hat{w}|}] \quad (1)$$

where $\hat{w}$ denotes the FFT of window w . Test features are normalised using training-set statistics to prevent data leakage.

C. Models

Statistical Threshold. Flags windows where standardised deviation or energy exceeds $\mu_{\text{train}} + 2\sigma_{\text{train}}$.

Isolation Forest. 200-tree ensemble; anomaly score $s = -f(\mathbf{x})$ (negated decision function). Contamination set to 0.15.

LSTM Autoencoder. Encoder LSTM (64 units, tanh, dropout 0.2) compresses to a bottleneck; decoder LSTM reconstructs the sequence. Per-window MSE is the anomaly score; threshold at $\mu_e + 2\sigma_e$.

D. Evaluation Protocols

Protocol A (Pooled). Random 70/30 split (858 train, 368 test windows).

Protocol B (Cross-subject). Subjects 1–21 train; subjects 22–30 form the held-out test set (759 train, 467 test windows). This tests generalisation to entirely unseen individuals—a prerequisite for clinical deployment.

E. SHAP Attribution

TreeSHAP [11] is applied to the fitted Isolation Forest to compute per-feature SHAP values. We report global mean |SHAP| and per-anomaly-type mean |SHAP| to identify which features activate for each clinical condition.

V. RESULTS

A. Protocol A: Pooled Evaluation

Table II summarises results. The LSTM Autoencoder achieves perfect precision (1.000) with MCC=0.711, meaning every window it flags is a true anomaly. It misses 45% of anomalies, but in a clinical context a system with zero false alarms is far more likely to be trusted and acted upon than one with high recall but an 87% false positive rate.

The Isolation Forest, despite a competitive AUC-ROC (0.710), flags 276 of 313 normal windows as anomalous in binary evaluation. This is an artefact of the fixed contamination parameter (0.15): the model is forced to label exactly 15% of test windows as anomalies regardless of score distribution. Its continuous score ranking is reasonable (AUC 0.710), but the operating point is poorly calibrated.

TABLE II
PROTOCOL A RESULTS — POOLED 70/30 SPLIT (368 TEST WINDOWS, 55 ANOMALIES)

Model	AUC-ROC	AUC-PR	F1	MCC	Prec	Rec
Threshold	0.718	0.748	0.000	0.000	0.000	0.000
Isolation Forest	0.710	0.748	0.211	−0.175	0.124	0.709
LSTM AE	0.722	0.749	0.706	0.711	1.000	0.545

B. Protocol B: Cross-Subject Generalisation

Table III reports cross-subject results. All models achieve *higher* AUC under Protocol B than Protocol A ($\Delta \approx +0.08$ across models). We attribute this to anomaly-type distribution variance: Protocol B’s random seed produced 23 fatigue-drift anomalies vs. 5 in Protocol A. Since fatigue drift is the hardest anomaly type (gradual slope, modest feature deviation), having more instances improves the models’ AUC by shifting the operating point on the ROC curve. This finding highlights that reported AUC can be sensitive to the composition of the anomaly set, motivating per-type evaluation (Section V-C).

The LSTM Autoencoder again achieves the strongest clinical operating point: perfect precision (1.000) with MCC=0.646, confirming its generalisation to unseen subjects.

TABLE III
PROTOCOL B RESULTS — CROSS-SUBJECT SPLIT (467 TEST WINDOWS, 70 ANOMALIES)

Model	AUC-ROC	AUC-PR	F1	MCC	Prec	Rec
Threshold	0.796	0.816	0.000	0.000	0.000	0.000
Isolation Forest	0.793	0.816	0.243	−0.038	0.144	0.786
LSTM AE	0.810	0.818	0.627	0.646	1.000	0.457

C. Per-Anomaly-Type Detection

Fig. ?? shows AUC-ROC per clinical anomaly type. Sensor noise and near-fall events are detected with $\text{AUC} \geq 0.93$ by all models—their large spectral variance and peak amplitude respectively produce strong feature responses. Compensatory asymmetry (amplitude attenuation) is the easiest type for

the threshold model (AUC 1.0) because attenuation directly reduces the energy feature below the training distribution. Progressive fatigue drift achieves $AUC \approx 0.91$ (Threshold), 0.91 (IF), and 0.91 (LSTM), confirming it as the hardest anomaly in absolute terms—but all models still detect it well above chance.

D. SHAP Attribution

Global SHAP analysis (Fig. ??) identifies `std_fft_magnitude` and `mean_fft_magnitude` as the two features with highest mean $|SHAP|$, dominating all time-domain statistics by a factor of 5–10 $\times$. This demonstrates that the Isolation Forest’s anomaly scores are driven almost entirely by spectral properties of the signal—not by raw amplitude statistics—a finding with direct implications for feature selection in future work.

The per-type SHAP plot (Fig. ??) reveals mechanistically distinct signatures:

- **Fatigue drift:** dominated by `mean_fft_magnitude` (the ramp shifts the DC component of the spectrum).
- **Sensor noise:** equal contributions from both FFT features and a notable `max_amplitude` contribution (broad-band noise raises the spectral floor and peak values).
- **Compensatory asymmetry:** uniquely shows elevated `std_fft_magnitude` with negligible `max_amplitude` (attenuation compresses amplitude range, increasing relative spectral variance).
- **Near-fall events:** highest `mean_fft_magnitude` and the largest `max_amplitude` contribution of any type (the 5-sample spike creates a sharp harmonic in the FFT and an extreme peak value).

These signatures are clinically interpretable and could guide the design of type-specific detectors in future work.

VI. DISCUSSION

A. Clinical Operating Point

The standard practice of reporting AUC-ROC obscures an important distinction for clinical deployment: a model with $AUC=0.71$ and perfect precision (LSTM) is fundamentally more deployable than one with $AUC=0.71$ and 87% false positive rate (Isolation Forest). Clinicians who receive false alarms rapidly lose trust in monitoring systems [13]. We recommend reporting MCC and precision-recall operating points alongside AUC for clinical anomaly detection work.

B. Generalisation Paradox

The counter-intuitive improvement from Protocol A to B (higher AUC on unseen subjects) is a warning about evaluation design. Anomaly-type composition—not just anomaly rate—substantially affects reported AUC. Future benchmarks should stratify results by anomaly type and report variance across random seeds.

C. Spectral Dominance

The SHAP finding that spectral features dominate time-domain statistics suggests that future feature sets for gait anomaly detection should prioritise FFT-derived representations. This aligns with prior work showing that gait has a characteristic spectral signature (cadence peak, harmonics) that is disrupted by most pathological conditions [2].

D. Limitations

The anomalies are synthetic. Validation against real pathological gait recordings is necessary. The dataset is limited to waist-mounted sensors; wrist or ankle generalisation requires additional study. The LSTM was trained for only 15 epochs; longer training may improve recall.

VII. CONCLUSION

We presented a clinically grounded, reproducible framework for unsupervised gait anomaly detection from IMU data, evaluated under both pooled and cross-subject protocols. The LSTM Autoencoder achieves the best clinical operating point—zero false alarms, $MCC \approx 0.65$ – 0.71 across both protocols—and generalises to unseen subjects without retraining. SHAP attribution reveals that spectral energy features drive detection across all anomaly types, with per-type signatures that are mechanistically interpretable. A counterintuitive improvement under cross-subject evaluation motivates careful attention to anomaly-type composition in benchmark design. Future work will validate the framework on real pathological gait datasets and explore longer-horizon temporal modelling for progressive anomaly types.

REFERENCES

- [1] World Health Organization, “Falls,” *WHO Fact Sheet*, 2021. [Online]. Available: <https://www.who.int/news-room/fact-sheets/detail/falls>
- [2] W. Tao, T. Liu, R. Zheng, and H. Feng, “Gait analysis using wearable sensors,” *Sensors*, vol. 12, no. 2, pp. 2255–2283, 2012.
- [3] J. Wang *et al.*, “Deep learning for sensor-based activity recognition: A survey,” *Pattern Recognit. Lett.*, vol. 119, pp. 3–11, 2019.
- [4] B. E. Maki, “Gait changes in older adults: predictors of falls or indicators of fear?” *J. Am. Geriatr. Soc.*, vol. 45, no. 3, pp. 313–320, 1997.
- [5] S. Mazilu *et al.*, “Online detection of freezing of gait with smartphones and machine learning,” in *Proc. 6th Int. Conf. Pervasive Comput. Technol. Healthcare*, 2012.
- [6] K. K. Patterson *et al.*, “Gait asymmetry in community-ambulating stroke survivors,” *Arch. Phys. Med. Rehabil.*, vol. 89, no. 2, pp. 304–310, 2008.
- [7] A. K. Bourke *et al.*, “Evaluation of a threshold-based tri-axial accelerometer fall detection algorithm,” *Gait & Posture*, vol. 26, no. 2, pp. 194–199, 2007.
- [8] G. S. Faber *et al.*, “Estimating 3D L5/S1 moments and ground reaction forces during trunk bending using a full-body ambulatory inertial motion capture system,” *J. Biomech.*, vol. 74, pp. 239–245, 2018.
- [9] F. T. Liu, K. M. Ting, and Z.-H. Zhou, “Isolation forest,” in *Proc. IEEE Int. Conf. Data Mining (ICDM)*, 2008, pp. 413–422.
- [10] P. Malhotra *et al.*, “LSTM-based encoder-decoder for multi-sensor anomaly detection,” *ICML Anomaly Detection Workshop*, 2016.
- [11] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Proc. NeurIPS*, 2017.
- [12] D. Anguita *et al.*, “A public domain dataset for human activity recognition using smartphones,” in *Proc. ESANN*, 2013.
- [13] M. Cvach, “Monitor alarm fatigue: an integrative review,” *Biomed. Instrum. Technol.*, vol. 46, no. 4, pp. 268–277, 2012.